

Name Kapil Uniyara ML Lab Work

Objective- The goal is to make some predictive models on a FHS dataset, and reviewing some exploratory and modelling techniques.

```
In [30]: #Reading the file
import pandas as pd
data = pd.read_csv("Kapil Uniyara - framingham - Kapil Uniyara - framingham.csv")
```

```
In [31]: #lets to undersand the data perform Some initial exploration
data.head()
```

Out[31]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diab
0	1	39	4.0	0	0.0	0.0	0	0	
1	0	46	2.0	0	0.0	0.0	0	0	
2	1	48	1.0	1	20.0	0.0	0	0	
3	0	61	3.0	1	30.0	0.0	0	1	
4	0	46	3.0	1	23.0	0.0	0	0	

```
In [32]: data.describe()
```

Out[32]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalent
count	4240.000000	4240.000000	4135.000000	4240.000000	4211.000000	4187.000000	4240.000000
mean	0.429245	49.580189	1.979444	0.494104	9.005937	0.029615	0.000000
std	0.495027	8.572942	1.019791	0.500024	11.922462	0.169544	0.000000
min	0.000000	32.000000	1.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	42.000000	1.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	49.000000	2.000000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	56.000000	3.000000	1.000000	20.000000	0.000000	0.000000
max	1.000000	70.000000	4.000000	1.000000	70.000000	1.000000	1.000000

```
In [33]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4240 entries, 0 to 4239
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4240 non-null   int64
1   age                   4240 non-null   int64
2   education             4135 non-null   float64
3   currentSmoker         4240 non-null   int64
4   cigsPerDay            4211 non-null   float64
5   BPMeds                4187 non-null   float64
6   prevalentStroke       4240 non-null   int64
7   prevalentHyp          4240 non-null   int64
8   diabetes              4240 non-null   int64
9   totChol               4190 non-null   float64
10  sysBP                 4240 non-null   float64
11  diaBP                 4240 non-null   float64
12  BMI                   4221 non-null   float64
13  heartRate             4239 non-null   float64
14  glucose               3852 non-null   float64
15  TenYearCHD            4240 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 530.1 KB
```

```
In [34]: data.isnull().sum()
```

```
Out[34]: male                0
age                0
education          105
currentSmoker      0
cigsPerDay         29
BPMeds             53
prevalentStroke    0
prevalentHyp       0
diabetes           0
totChol            50
sysBP              0
diaBP              0
BMI                19
heartRate          1
glucose            388
TenYearCHD         0
dtype: int64
```

Now we have some missing values, So lets make some visualization of it

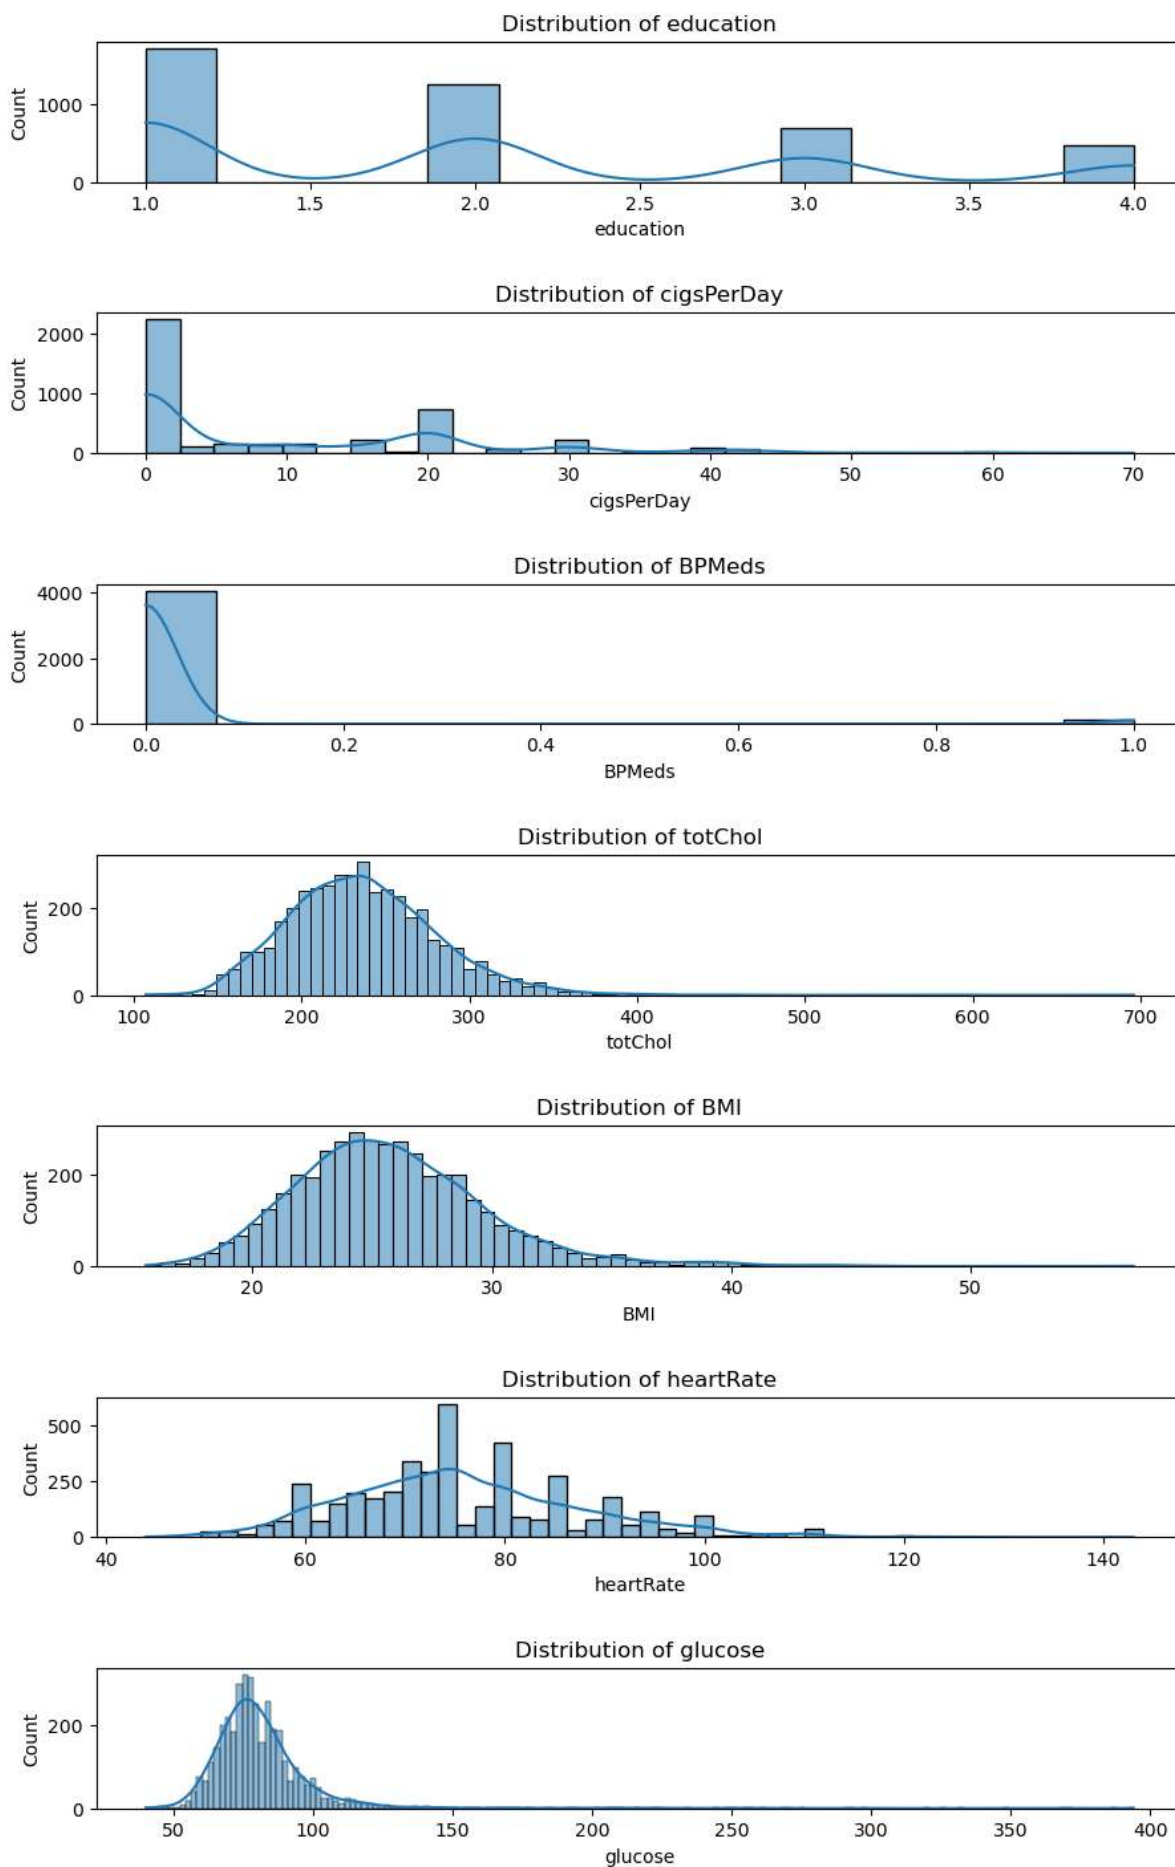
```
In [35]: import matplotlib.pyplot as plt
import seaborn as sns

cols_with_missing_values = ['education', 'cigsPerDay', 'BPMeds', 'totChol', 'BMI',

# Plotting the distributions of columns with missing values
fig, axes = plt.subplots(len(cols_with_missing_values), 1, figsize=(10, 15))
fig.tight_layout(pad=5.0)

for i, col in enumerate(cols_with_missing_values):
    sns.histplot(data[col], kde=True, ax=axes[i])
    axes[i].set_title(f'Distribution of {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Count')
```

```
plt.show()
```



Now lets fill missng values

```
In [36]: from sklearn.impute import SimpleImputer
# For normally distributed or symmetric continuous variables, we'll use the mean
# For skewed continuous variables and categorical variables, we'll use the median

# Imputer for mean
mean_imputer = SimpleImputer(strategy='mean')

# Columns to impute with mean (normally distributed/symmetric)
mean_cols = ['heartRate']

# Impute the columns with mean
data[mean_cols] = mean_imputer.fit_transform(data[mean_cols])
```

```
In [37]: # Imputer for median
median_imputer = SimpleImputer(strategy='median')

# Columns to impute with median (skewed distributions or categorical)
median_cols = ['education', 'cigsPerDay', 'BPMeds', 'totChol', 'BMI', 'glucose']

# Impute the columns with median
data[median_cols] = median_imputer.fit_transform(data[median_cols])
```

```
In [38]: # Checking if there are any missing values left
data.isnull().sum()
```

```
Out[38]: male                0
age                0
education          0
currentSmoker      0
cigsPerDay         0
BPMeds             0
prevalentStroke    0
prevalentHyp       0
diabetes           0
totChol            0
sysBP              0
diaBP              0
BMI                0
heartRate          0
glucose            0
TenYearCHD         0
dtype: int64
```

```
In [39]: data.describe()
```

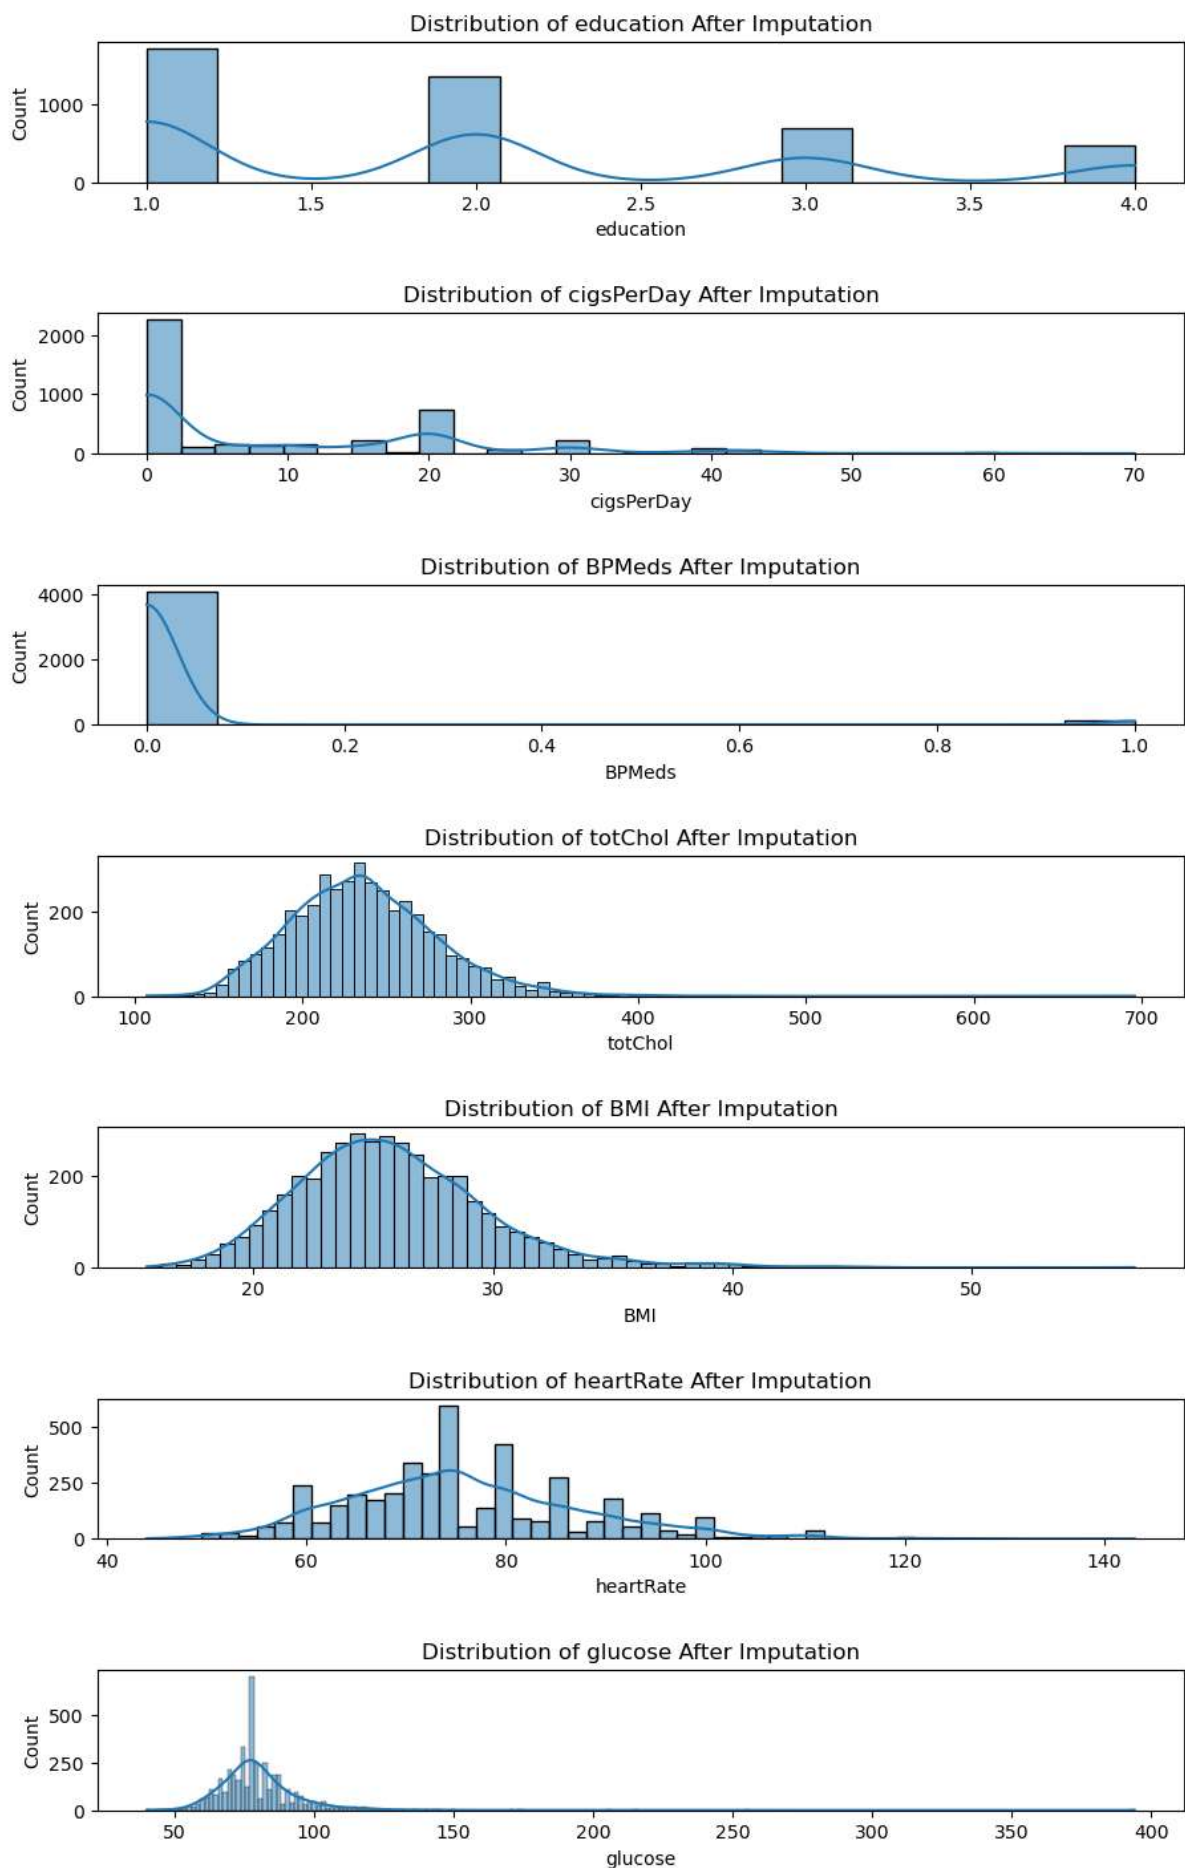
Out[39]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalent
count	4240.000000	4240.000000	4240.000000	4240.000000	4240.000000	4240.000000	4240.000000
mean	0.429245	49.580189	1.979953	0.494104	8.944340	0.029245	0.000000
std	0.495027	8.572942	1.007087	0.500024	11.904777	0.168513	0.000000
min	0.000000	32.000000	1.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	42.000000	1.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	49.000000	2.000000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	56.000000	3.000000	1.000000	20.000000	0.000000	0.000000
max	1.000000	70.000000	4.000000	1.000000	70.000000	1.000000	1.000000

```
In [40]: # Plotting the distributions after imputation
fig, axes = plt.subplots(len(cols_with_missing_values), 1, figsize=(10, 15))
fig.tight_layout(pad=5.0)

for i, col in enumerate(cols_with_missing_values):
    sns.histplot(data[col], kde=True, ax=axes[i])
    axes[i].set_title(f'Distribution of {col} After Imputation')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Count')

plt.show()
```










Now lets see the outliers in the data


```
In [41]: # Selecting continuous variables for outlier detection
         continuous_cols = ['age', 'cigsPerDay', 'totChol', 'sysBP', 'diaBP', 'BMI', 'heartRate', 'glucose']
```

```
# Plotting the boxplots for continuous variables
fig, axes = plt.subplots(len(continuous_cols), 1, figsize=(10, 20))
fig.tight_layout(pad=5.0)

for i, col in enumerate(continuous_cols):
    sns.boxplot(x=data[col], ax=axes[i])
    axes[i].set_title(f'Boxplot of {col}')
    axes[i].set_xlabel(col)

plt.show()
```





We will be using interquartile range (IQR) method to handle outliers

```
In [42]: def handle_outliers_with_IQR(df, column):
          """ Handle outliers in a dataframe column using the IQR method.
              Values outside 1.5 * IQR from the Q1 and Q3 quartiles are considered outliers
          """
          Q1 = df[column].quantile(0.25)
          Q3 = df[column].quantile(0.75)
          IQR = Q3 - Q1

          lower_bound = Q1 - 1.5 * IQR
          upper_bound = Q3 + 1.5 * IQR


          df[column] = df[column].clip(lower=lower_bound, upper=upper_bound)
```



```
In [43]: # Applying the IQR method to each continuous column
          for col in continuous_cols:
              handle_outliers_with_IQR(data, col)

          # Plotting the boxplots for continuous variables after handling outliers
          fig, axes = plt.subplots(len(continuous_cols), 1, figsize=(10, 20))
          fig.tight_layout(pad=5.0)

          for i, col in enumerate(continuous_cols):
              sns.boxplot(x=data[col], ax=axes[i])
              axes[i].set_title(f'Boxplot of {col} After Handling Outliers')
              axes[i].set_xlabel(col)

          plt.show()
```






Now lets check the data is imbalanced or not

```
In [44]: # Checking if the data is imbalanced
target_column = 'TenYearCHD'
class_counts = data[target_column].value_counts()

# Plotting the class distribution
plt.figure(figsize=(8, 5))
sns.barplot(x=class_counts.index, y=class_counts.values)
plt.title('Class Distribution of the Target Variable (TenYearCHD)')
plt.xlabel('Class')
plt.ylabel('Count')
plt.show()
```



```
In [45]: class_counts
```

```
Out[45]: TenYearCHD
0      3596
1       644
Name: count, dtype: int64
```

We will be using oversampling to balance the dataset

```
In [46]: from sklearn.utils import resample
```

```
# Separate the minority and majority classes
data_majority = data[data[target_column] == 0]
data_minority = data[data[target_column] == 1]
```

```
In [47]: # Upsample minority class
data_minority_upsampled = resample(data_minority,
                                   replace=True,      # sample with replacement
                                   n_samples=len(data_majority), # to match majority class
                                   random_state=123) # reproducible results

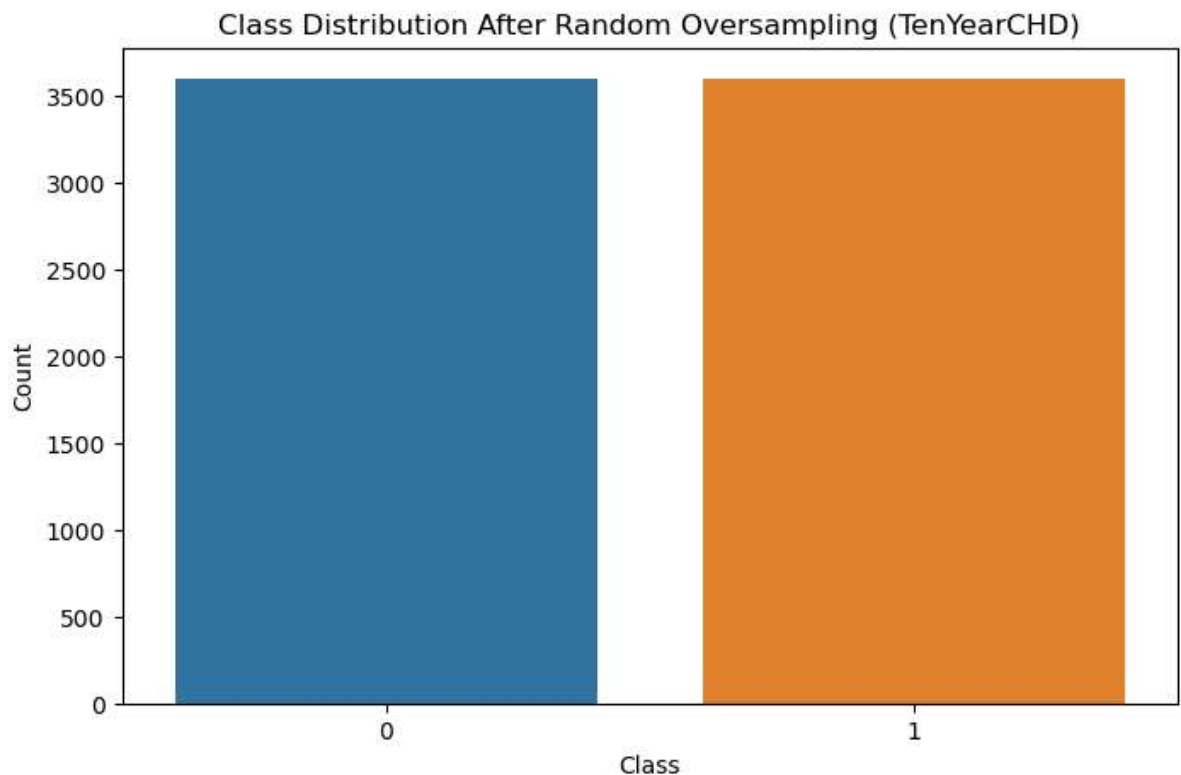
# Combine majority class with upsampled minority class
data_upsampled = pd.concat([data_majority, data_minority_upsampled])

# Display new class counts
upsampled_class_counts = data_upsampled[target_column].value_counts()
```

```
In [48]: upsampled_class_counts
```

```
Out[48]: TenYearCHD
0      3596
1      3596
Name: count, dtype: int64
```

```
In [49]: # Plotting the class distribution after resampling
plt.figure(figsize=(8, 5))
sns.barplot(x=upsampled_class_counts.index, y=upsampled_class_counts.values)
plt.title('Class Distribution After Random Oversampling (TenYearCHD)')
plt.xlabel('Class')
plt.ylabel('Count')
plt.show()
```



Now lets make predictions

```
In [50]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score
from sklearn.utils import resample

# Splitting the dataset into features and target variable
X = data_upsampled.drop('TenYearCHD', axis=1)
y = data_upsampled['TenYearCHD']

# Splitting the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [51]: # Initialize the Random Forest classifier
rf_classifier = RandomForestClassifier(random_state=42)

# Train the model
rf_classifier.fit(X_train, y_train)

# Make predictions on the test set
y_pred = rf_classifier.predict(X_test)
```

```
In [52]: accuracy_score(y_test, y_pred)
```

```
Out[52]: 0.9694232105628909
```

```
In [55]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.99	0.95	0.97	735
1	0.95	0.99	0.97	704
accuracy			0.97	1439
macro avg	0.97	0.97	0.97	1439
weighted avg	0.97	0.97	0.97	1439

Saving the pickle file to make webui for this problem

```
In [56]: import pickle

# Save the model to a pickle file
pickle_filename = 'random_forest_model.pkl'
with open(pickle_filename, 'wb') as file:
    pickle.dump(rf_classifier, file)

print(f"Model saved as {pickle_filename}")
```

Model saved as random_forest_model.pkl

```
In [ ]:
```